The effectiveness of moderating harmful online content

created by DALL-E 3,
prompt “The effectiveness of moderating harmful online content”

Philipp Schneider
Risk Analytics and Optimization @ EPFL
philipp.schneider@epfl.ch

Dr Marian-Andrei Rizoiu
Behavioral Data Science @ UTS
marian-andrei.rizoiu@uts.edu.au

Harmful content in numbers: Trends and statistics

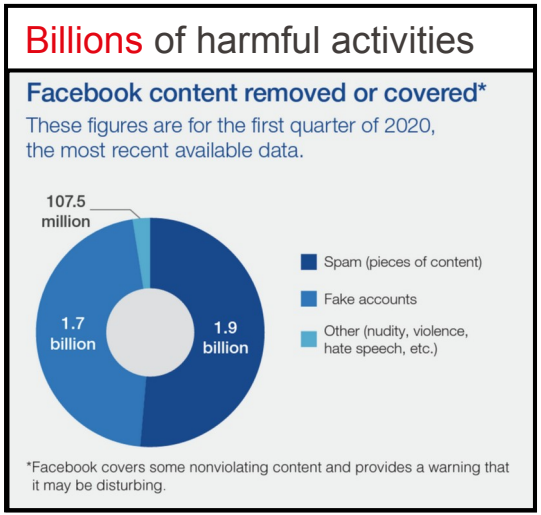

[1]

Harmful content in numbers: Trends and statistics

Billions of harmful activities

Facebook content removed or covered*

These figures are for the first quarter of 2020, the most recent available data.

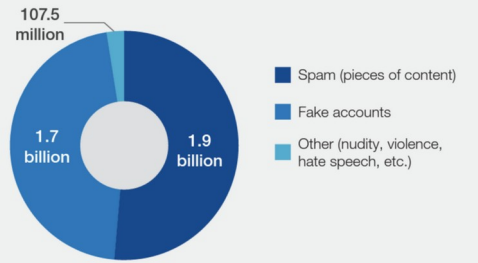

*Facebook covers some nonviolating content and provides a warning that it may be disturbing.

[1]

Diverse spectrum of harmful content categories

Facebook removals other than fake accounts and spam*

First quarter of 2020, in millions.

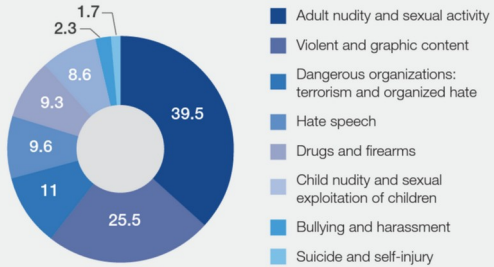

*Includes some content that is covered but not removed.

Harmful content in numbers: Trends and statistics

Billions of harmful activities

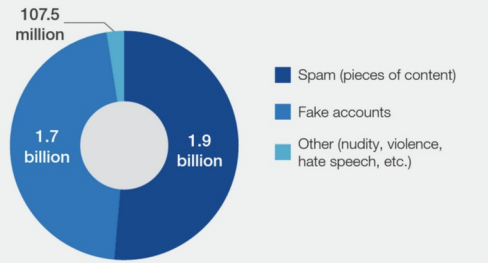

[1]

Diverse spectrum of harmful content categories

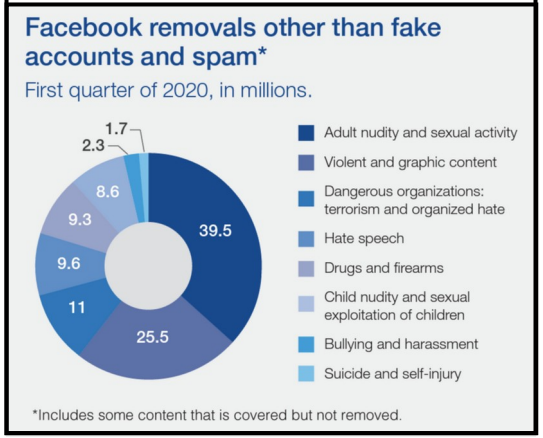

Algorithms are increasingly vital for detection

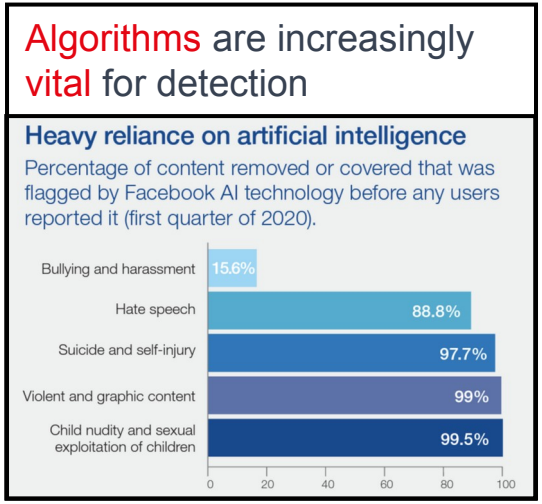

Harmful content in numbers: Trends and statistics

Billions of harmful activities

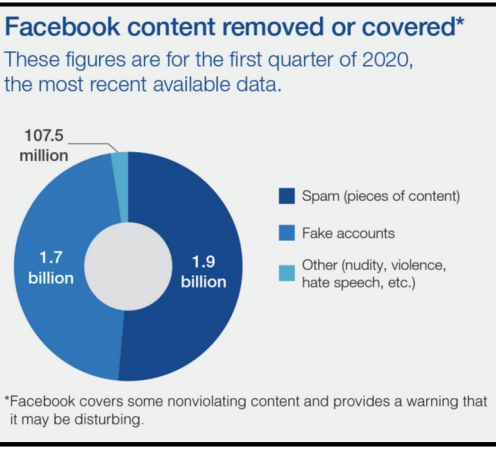

[1]

Diverse spectrum of harmful content categories

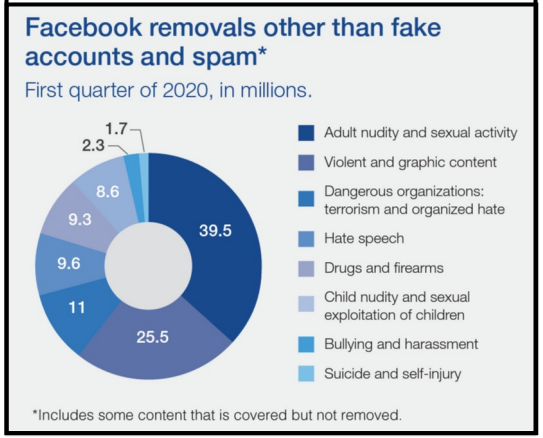

Algorithms are increasingly vital for detection

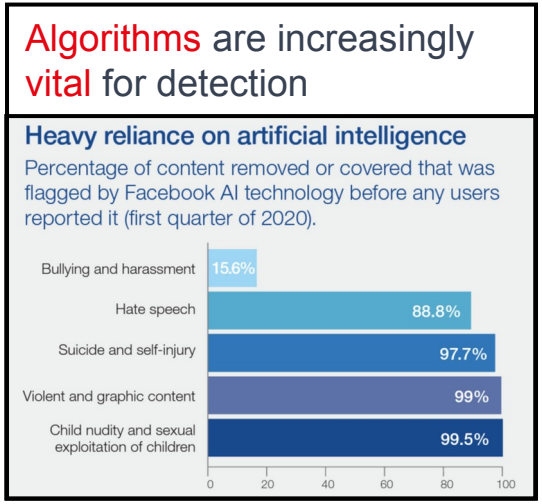

Harmful content

- **Misinformation:** Dissemination of false or inaccurate information without proper knowledge or verification
- **Disinformation:** Intentionally created with the aim of misleading and disseminating false information (subclass: Illegal content)

Society's perception of content moderation is evolving ...

Support for the U.S. government and tech companies restricting false information online has risen steadily in recent years

% of U.S. adults who say ...

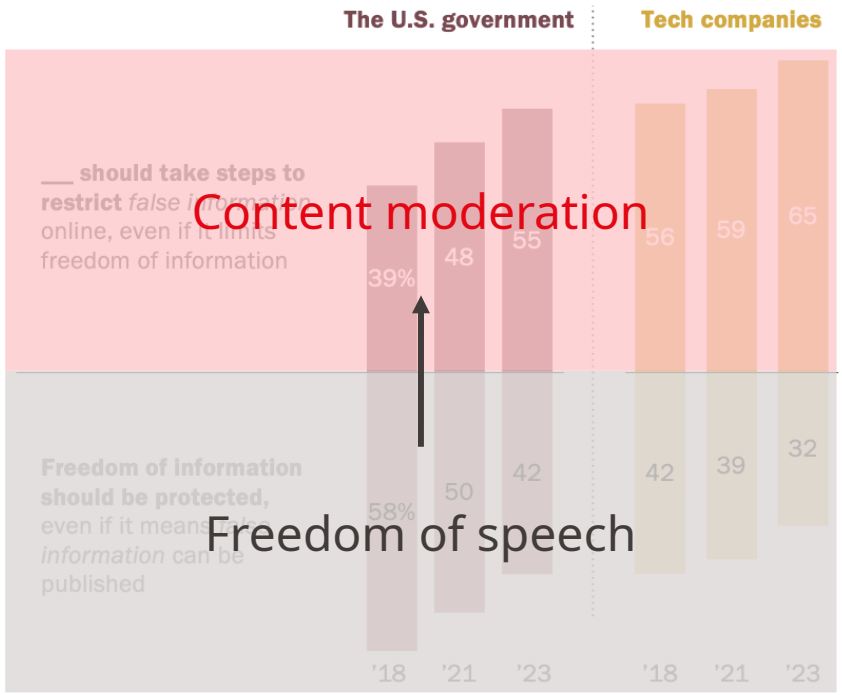

Note: Respondents who did not answer are not shown.
Source: Survey of U.S. adults conducted June 5-11, 2023.

PEW RESEARCH CENTER

[2]

Society's perception of content moderation is evolving ...

Support for the U.S. government and tech companies restricting false information online has risen steadily in recent years

% of U.S. adults who say ...

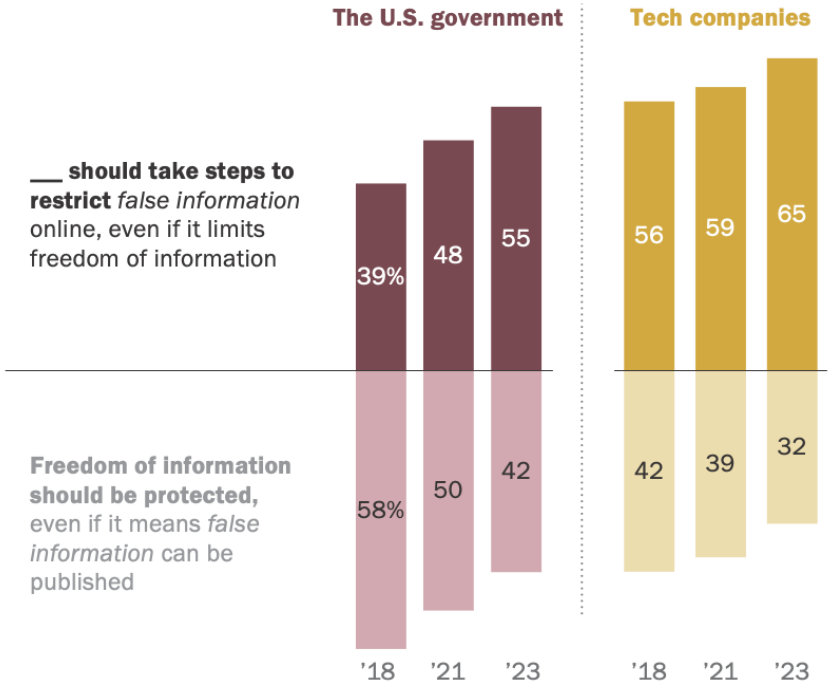

Note: Respondents who did not answer are not shown.
Source: Survey of U.S. adults conducted June 5-11, 2023.

PEW RESEARCH CENTER

[2]

Generative AI's impact on misinformation: Solution or problem?

OpenAI says AI tools can be effective in content moderation

Reuters

August 15, 2023 9:27 PM GMT+2 · Updated a month ago

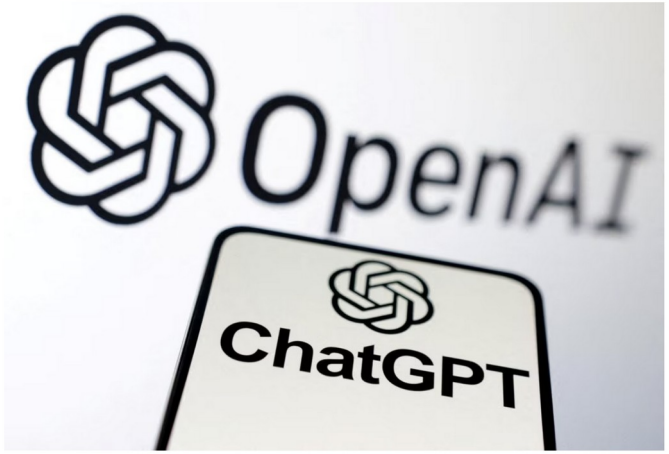

OpenAI and ChatGPT logos are seen in this illustration taken, February 3, 2023. REUTERS/Dado Ruvic/Illustration/File Photo *Acquire Licensing Rights*

[3]

Generative AI's impact on misinformation: Solution or problem?

OpenAI says AI tools can be effective in content moderation

Reuters

August 15, 2023 9:27 PM GMT+2 · Updated a month ago

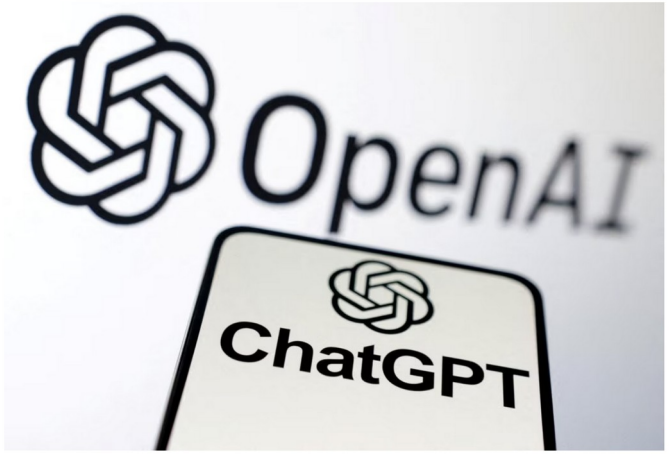

OpenAI and ChatGPT logos are seen in this illustration taken, February 3, 2023. REUTERS/Dado Ruvic/Illustration/File Photo *Acquire Licensing Rights*

[3]

Tech firms leverage generative AI to combat misinformation, driving workforce adaptations.

Tech layoffs shrink 'trust and safety' teams, raising fears of backsliding efforts to curb online abuse

"Fewer people means less work is being done in a lot of different spaces," said one of Twitter's remaining content moderation staffers.

The Twitter headquarters in San Francisco, on Dec. 8, 2022. Jeff Chiu / AP file

[4]

Humans may be more likely to believe disinformation generated by AI

The way AI models structure text may have something to do with it, according to the study authors.

By Rhiannon Williams

June 28, 2023

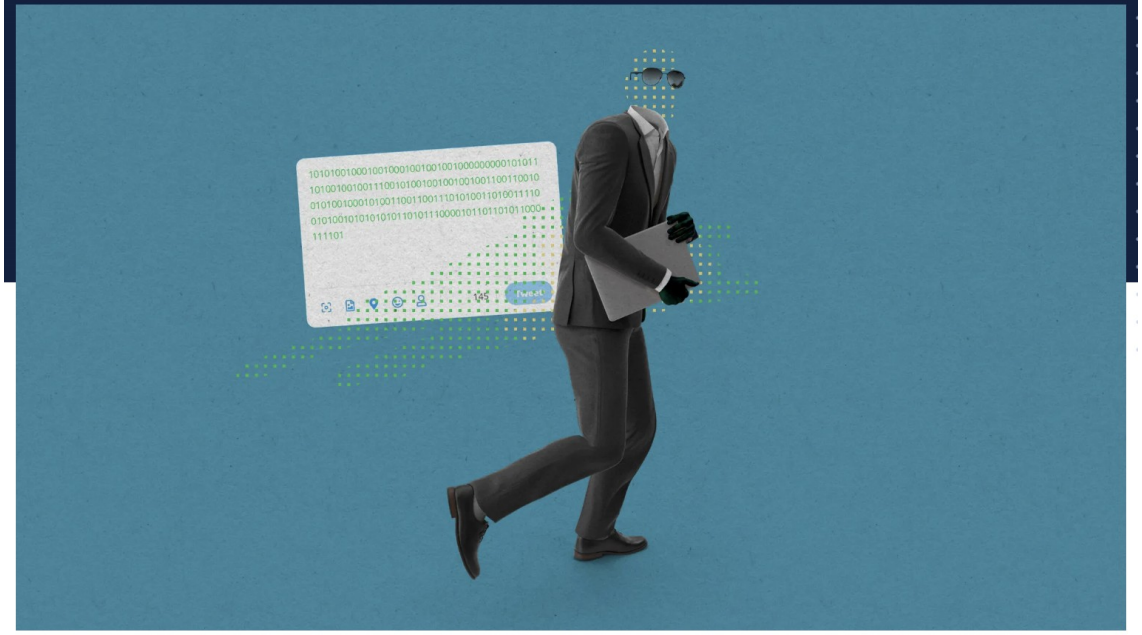

STEPHANIE ARNETT/MITTR | ENVATO

[5]

Generative AI's impact on misinformation: Solution or problem?

Researchers demonstrate that generative AI (GPT) is capable of generating more persuasive disinformation. [6]

EU's content removal strategy incorporates a human element

Digital Services Act – Content moderation

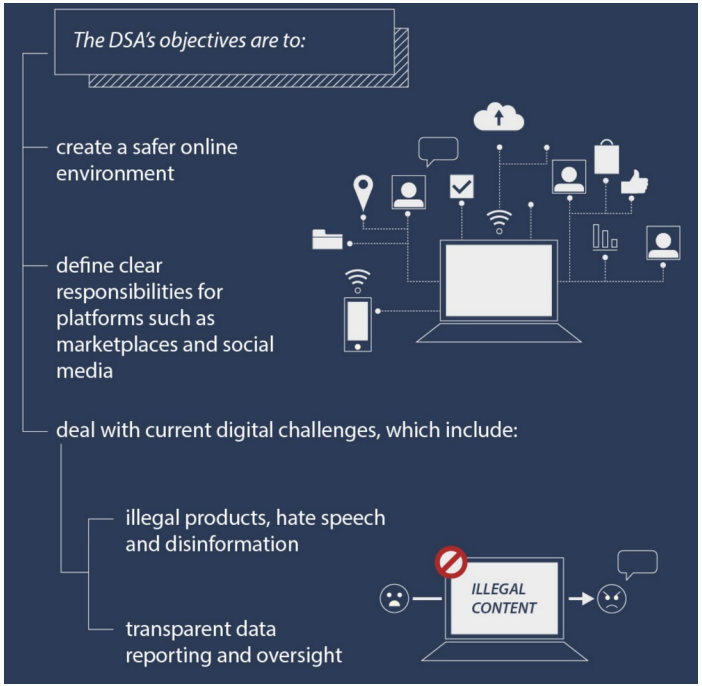

@DSA-Infographic

Trusted flagger mechanism [7]

"Trusted flagger" (officially appointed entity) reporting problematic content to platforms, who must then remove it within 24 hours.

Digital Services Act – Content moderation

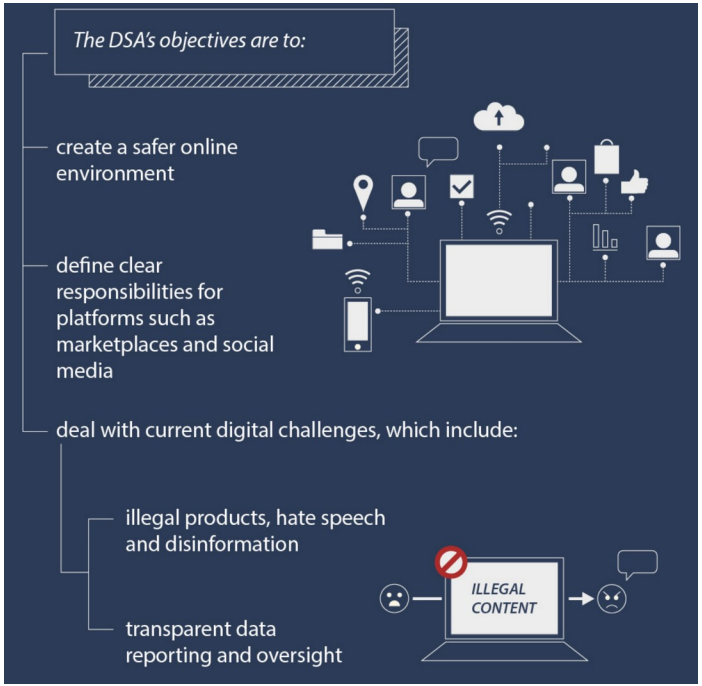

@DSA-Infographic

Trusted flagger mechanism [7]

"Trusted flagger" (officially appointed entity) reporting problematic content to platforms, who must then remove it within 24 hours.

Can human moderators ever really rein in harmful online content?

PNAS BRIEF REPORT | SOCIAL SCIENCES OPEN ACCESS

The effectiveness of moderating harmful online content

Philipp J. Schneider[a,1] and Marian-Andrei Rizoiu[b,1,2]

Edited by Margaret Levi, Stanford University, Sanford, CA; received May 2, 2023; accepted June 28, 2023

In 2022, the European Union introduced the Digital Services Act (DSA), a new legislation to report and moderate harmful content from online social networks. Trusted flaggers are mandated to identify harmful content, which platforms must remove within a set delay (currently 24 h). Here, we analyze the likely effectiveness of EU-mandated mechanisms for regulating highly viral online content with short half-lives. We deploy self-exciting point processes to determine the relationship between the regulated moderation delay and the likely harm reduction achieved. We find that harm reduction is achievable for the most harmful content, even for fast-paced platforms such as Twitter. Our method estimates moderation effectiveness for a given platform and provides a rule of thumb for selecting content for investigation and flagging, managing flaggers' workload.

content moderation | harmful content | harm reduction | stochastic modeling

Social media platforms are the new town squares (1)—dematerialized, digital, and unregulated town squares. In 2022, Elon Musk acquired Twitter with the stated goal of preserving free speech for the future. However, alongside free speech, harmful content disseminates and prospers in this unregulated space: mis- and disinformation that spreads faster than its debunking (2), social bots that infiltrate political processes (3), hate speech against women, immigrants, and minorities (4) or viral challenges that put teens' lives

Schneider, P. J., & Rizoiu, M.-A. (2023). The effectiveness of moderating harmful online content. Proceedings of the National Academy of Sciences, 120(34), 1–3. https://doi.org/10.1073/pnas.2307360120

PNAS BRIEF REPORT | SOCIAL SCIENCES

The effectiveness of moderating harmful online

Philipp J. Schneider[a,1] and Marian-Andrei Rizoiu[b,1,2]

Edited by Margaret Levi, Stanford University, Sanford, CA; received May 2, 2023; accepted June 28, 2023

In 2022, the European Union introduced the Digital Services Act (DSA), a new legislation to report and moderate harmful content from online social networks. Trusted flaggers are mandated to identify harmful content, which platforms must remove within a set delay (currently 24 h). Here, we analyze the likely effectiveness of EU-mandated mechanisms for regulating highly viral online content with short half-lives. We deploy self-exciting point processes to determine the relationship between the regulated moderation delay and the likely harm reduction achieved. We find that harm reduction is achievable for the most harmful content, even for fast-paced platforms such as Twitter. Our method estimates moderation effectiveness for a given platform and provides a rule of thumb for selecting content for investigation and flagging, managing flaggers' workload.

content moderation | harmful content | harm reduction | stochastic modeling

Social media platforms are the new town squares (1)—dematerialized, digital, and unregulated town squares. In 2022, Elon Musk acquired Twitter with the stated goal of preserving free speech for the future. However, alongside free speech, harmful content disseminates and prospers in this unregulated space: mis- and disinformation that spreads faster than its debunking (2), social bots that infiltrate political processes (3), hate speech against women, immigrants, and minorities (4) or viral challenges that put teens' lives

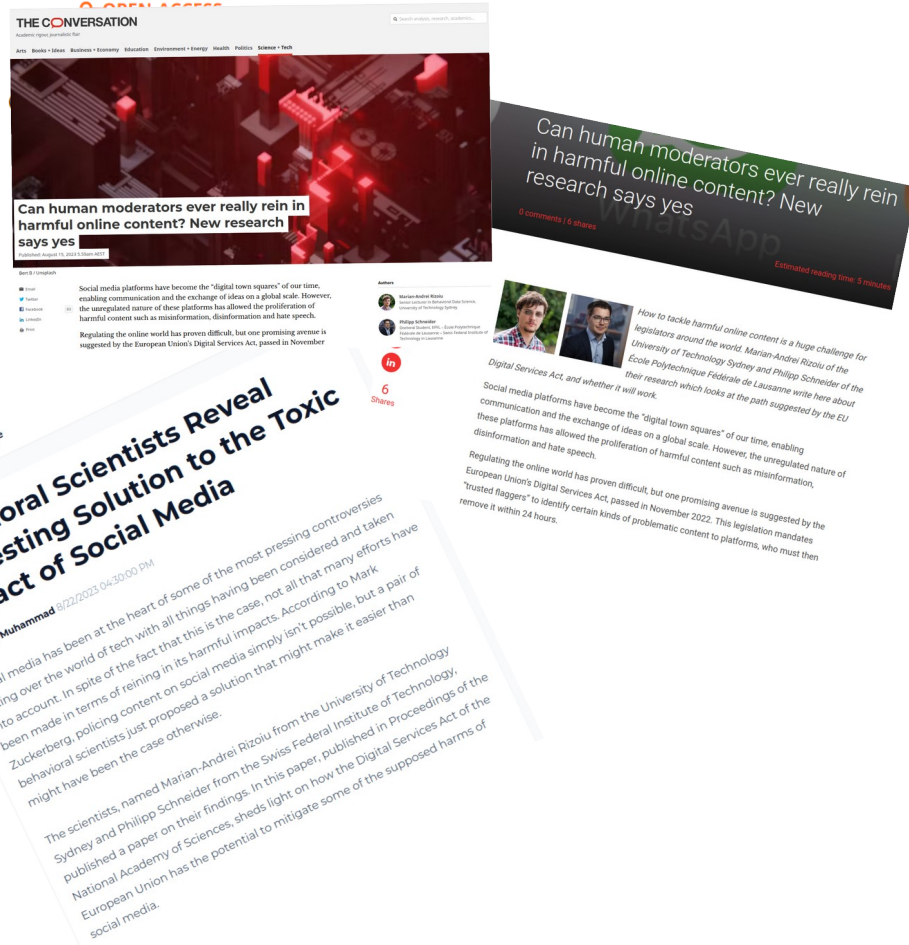

Content dynamics of harmful content

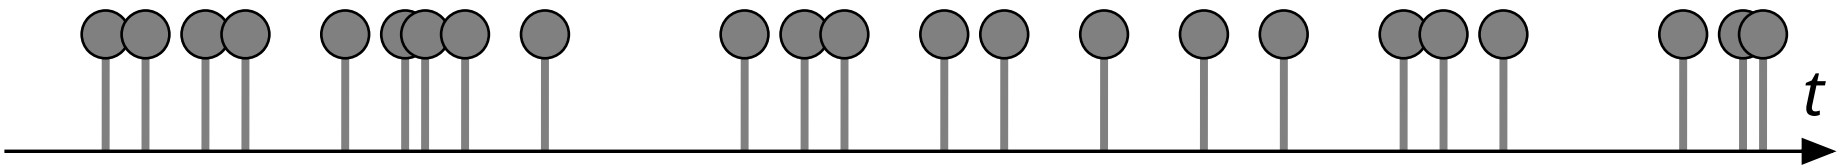

Posts

Content dynamics of harmful content

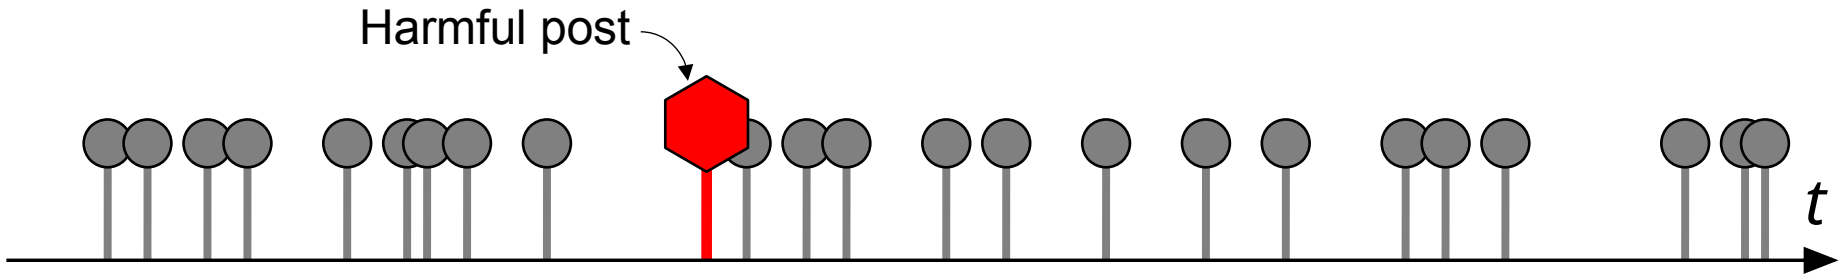

Posts

Content dynamics of harmful content

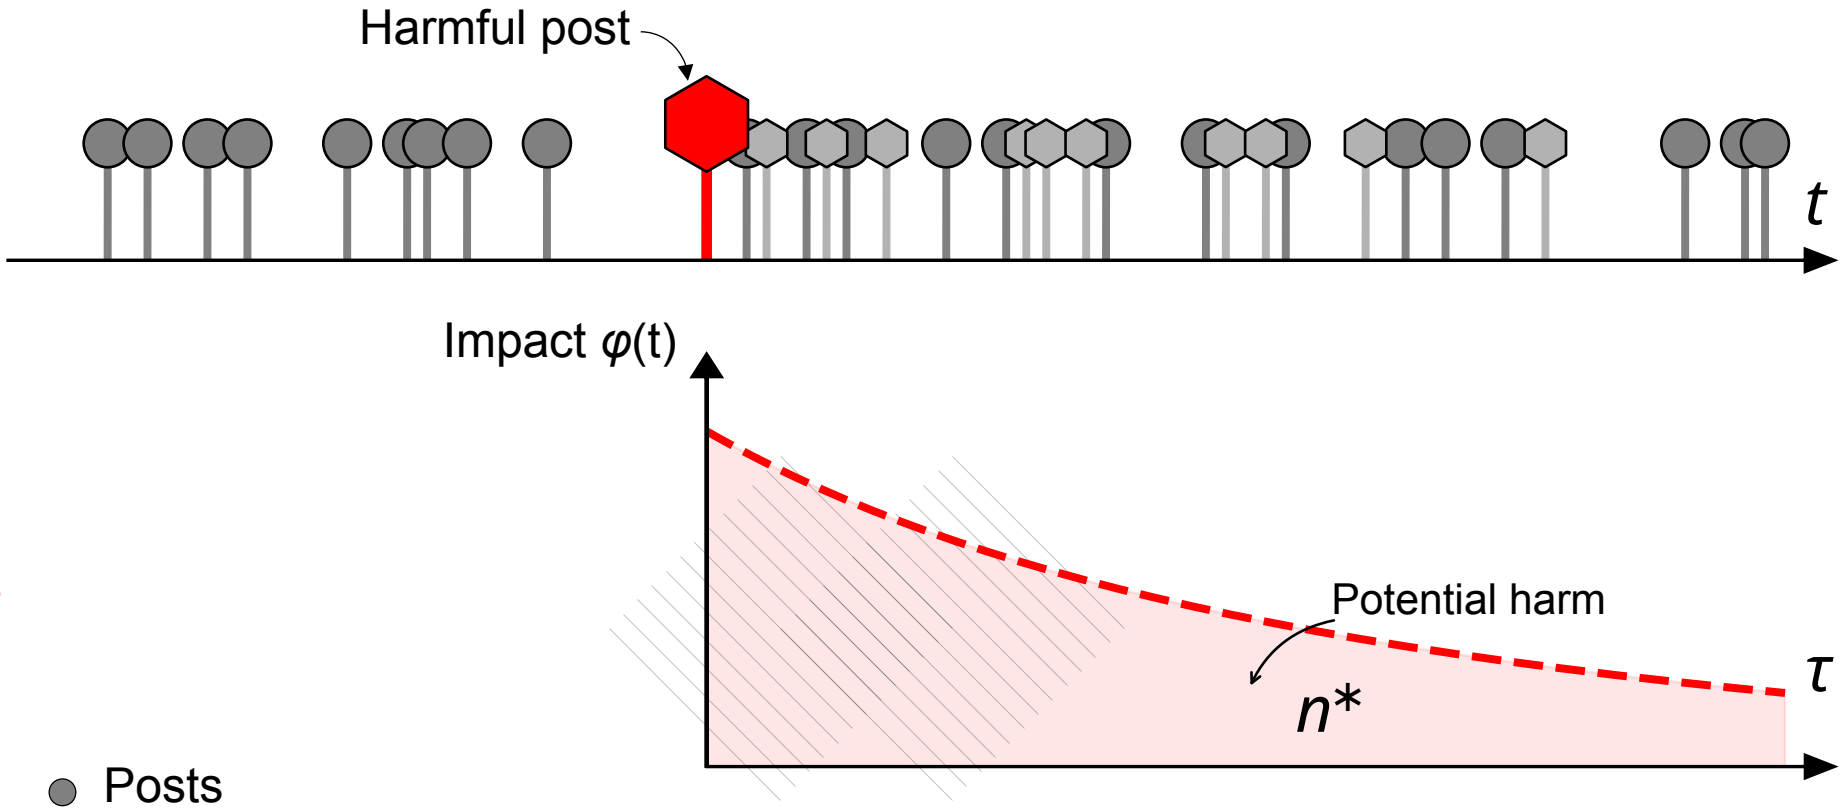

Content dynamics of harmful content

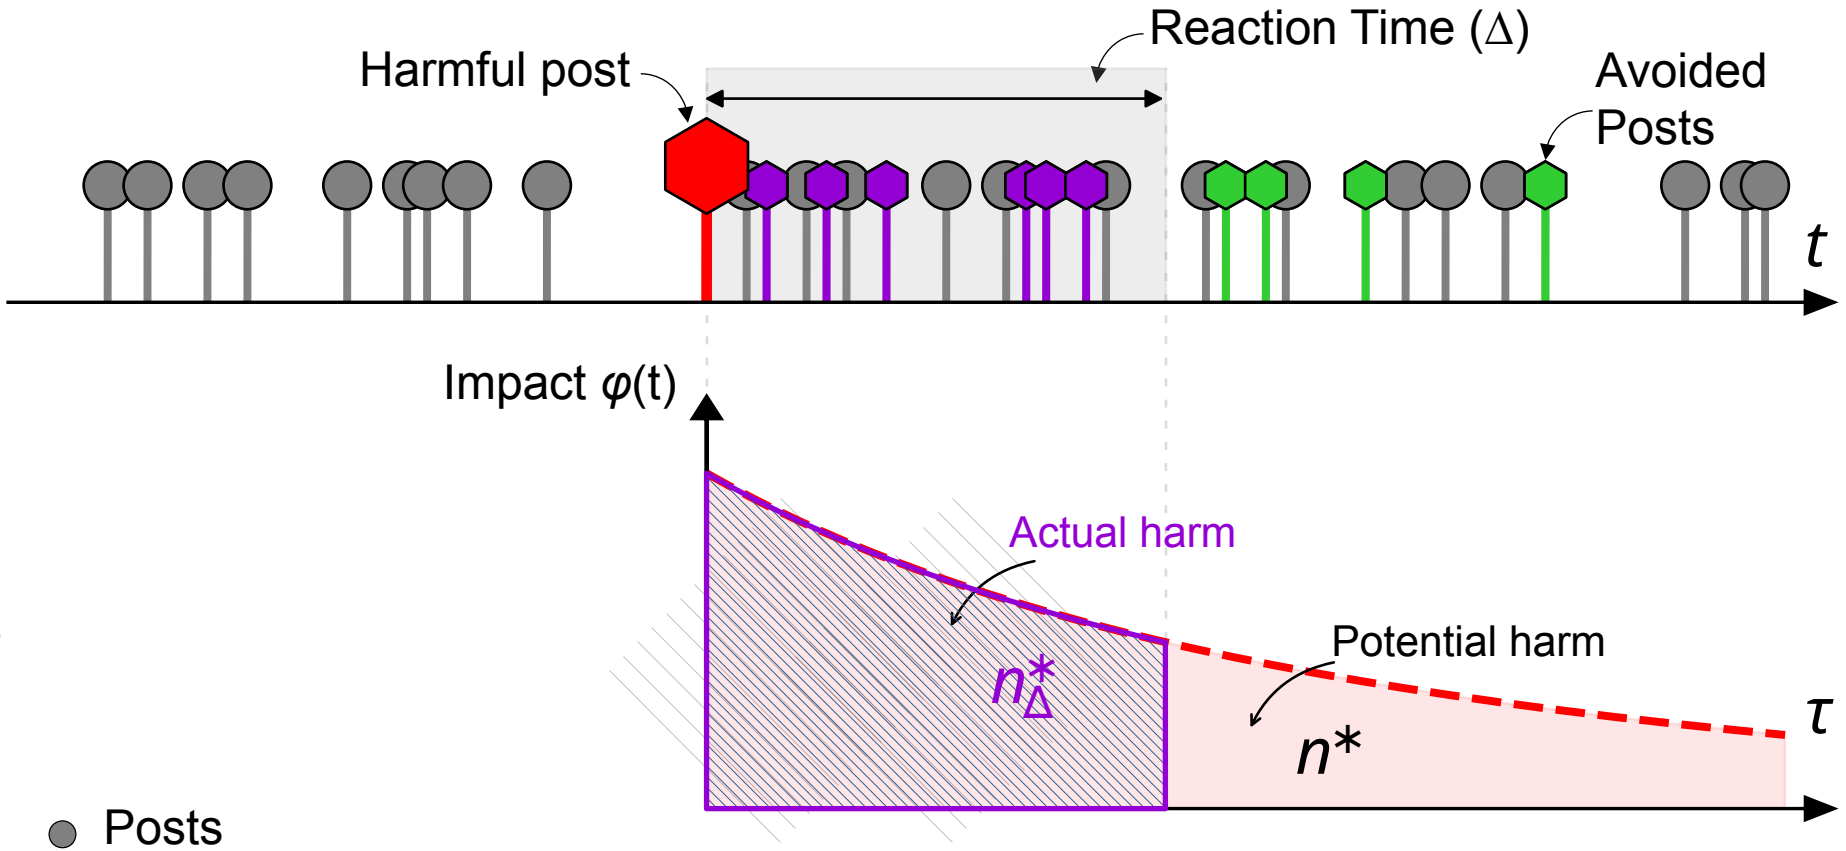

Real-world events occur in groups

Arrival Processes

Homogenous Poisson point process

$$\lambda(t) = \mu$$

Real-world events occur in groups

Arrival Processes

Homogenous Poisson point process

$$\lambda(t) = \mu$$

Self-exciting point process

Applications

Seismology

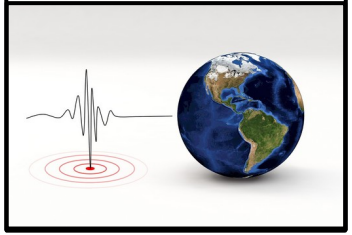

Finance

Social Media

Self-excitation in social media

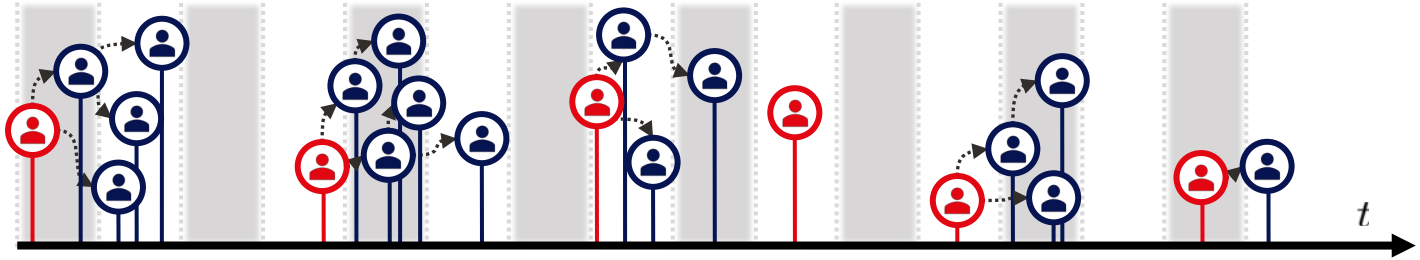

Self-excitation in social media

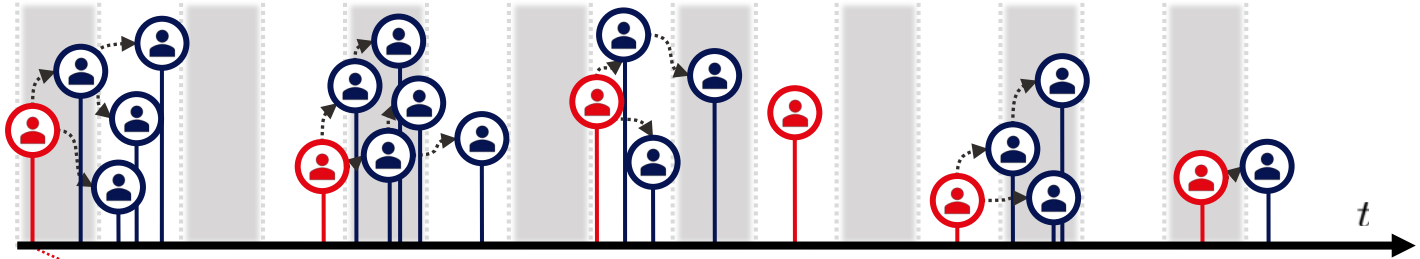

Self-excitation in social media

Original posts

Retweets / Likes

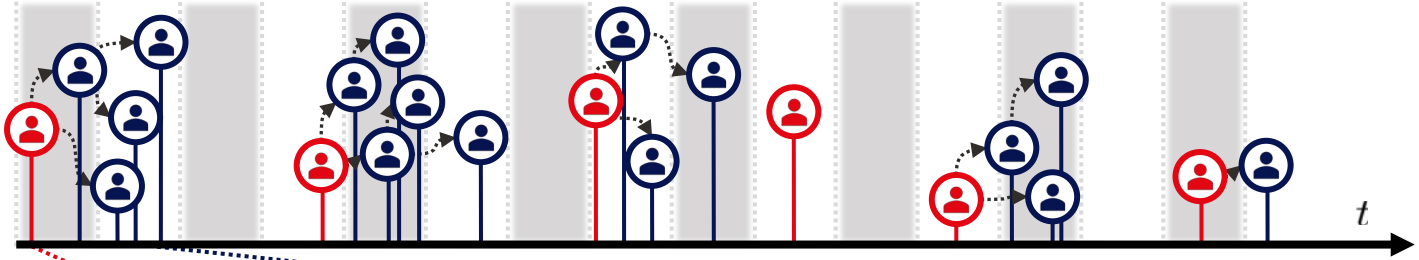

Self-excitation describes the clustering effect

Base intensity
(exogenous)

Intensity function $\lambda(t|\mathcal{H}_t) = \mu + \sum_{i:t_i<t} \phi(t - t_i)$

Event history

$\mathcal{H}_T = \{t_1, \ldots, t_K\} \subset (0, T]$

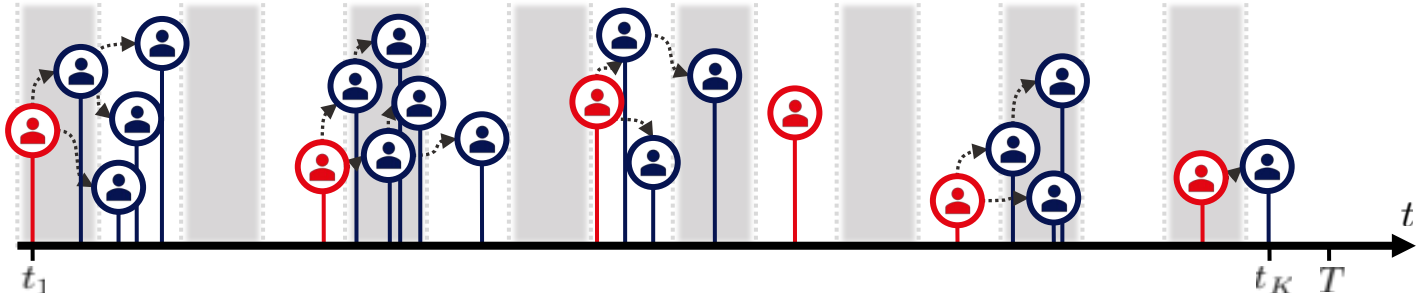

Self-excitation describes the clustering effect

Base intensity (exogenous)

Self-excitation (endogenous)

Intensity function

$$\lambda(t|\mathcal{H}_t) = \mu + \sum_{i:t_i<t} \phi(t - t_i)$$

Event history

$$\mathcal{H}_T = \{t_1, \ldots, t_K\} \subset (0, T]$$

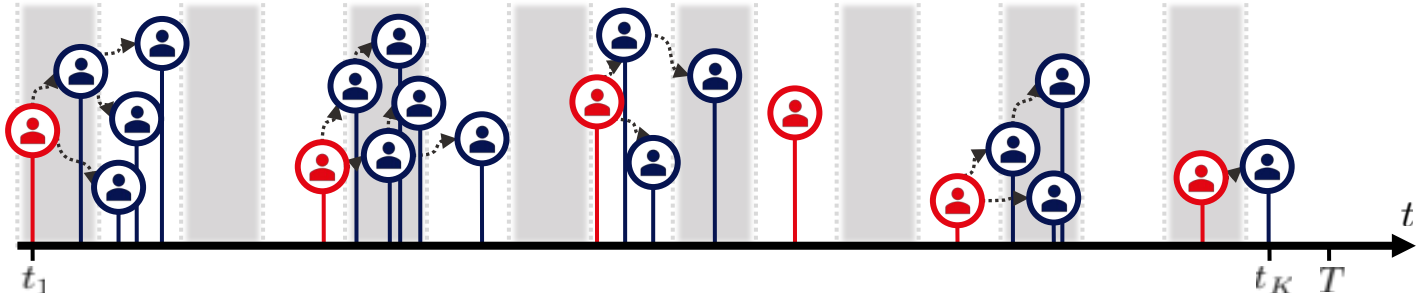

Self-excitation describes the clustering effect

Base intensity (exogenous)

Self-excitation (endogenous)

Intensity function

$$\lambda(t|\mathcal{H}_t) = \mu + \sum_{i:t_i<t} \phi(t - t_i)$$

Kernel function

Event history

$$\mathcal{H}_T = \{t_1, \ldots, t_K\} \subset (0, T]$$

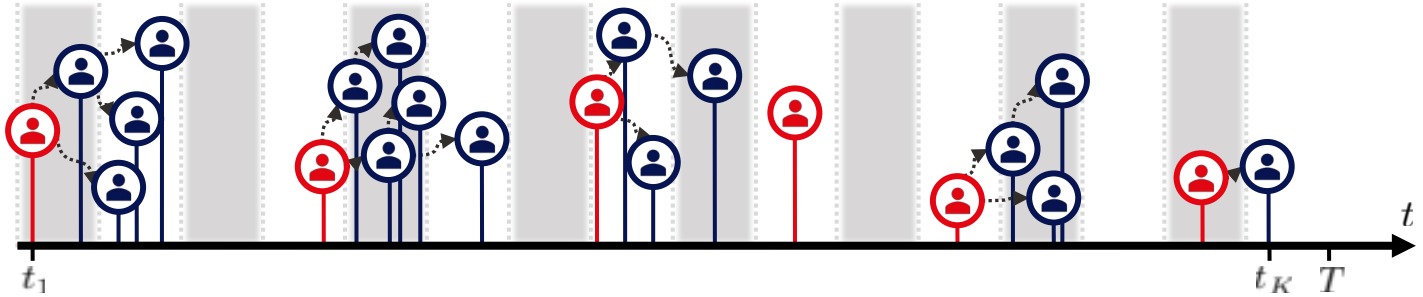

What are the key metrics of this study?

Potential harm: Number of harmful offspring the post generates

Content half-life: Amount of time required for half of all the post's offspring to be generated

What are the key metrics of this study?

Potential harm: Number of harmful offspring the post generates

Content half-life: Amount of time required for half of all the post's offspring to be generated

- Twitter – 24 min
- Facebook – 105 min
- Instagram – 20 h
- LinkedIn – 24 h
- YouTube – 8.8 d
- Pinterest – 3.75 mo

[8]

What are the key metrics of this study?

Potential harm: Number of harmful offspring the post generates

Content half-life: Amount of time required for half of all the post's offspring to be generated

- Twitter – 24 min
- Facebook – 105 min
- Instagram – 20 h
- LinkedIn – 24 h
- YouTube – 8.8 d
- Pinterest – 3.75 mo

[8]

News around terrorist attack

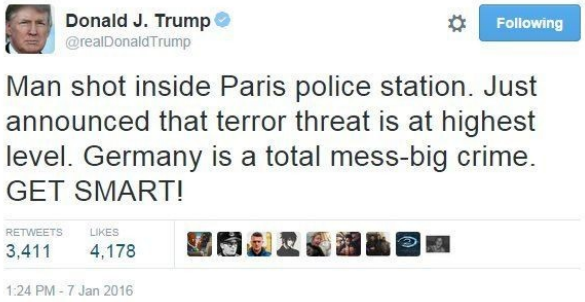

@BBC

- High potential harm / virality
- Short content half-life

What are the key metrics of this study?

Potential harm: Number of harmful offspring the post generates

Content half-life: Amount of time required for half of all the post's offspring to be generated

- Twitter – 24 min
- Facebook – 105 min
- Instagram – 20 h
- LinkedIn – 24 h
- YouTube – 8.8 d
- Pinterest – 3.75 mo

[8]

News around terrorist attack

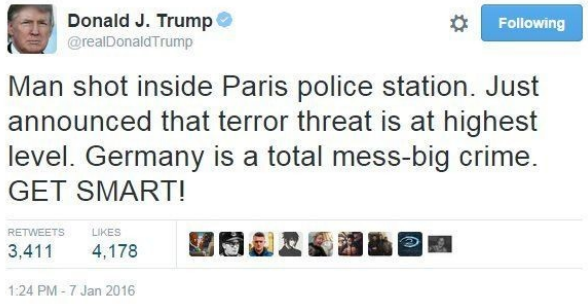

@BBC

- High potential harm / virality
- Short content half-life

Anti-vaccine conspiracies

@Vox

- Low potential harm / virality (before COVID-19)
- Long content half-life

What is the reaction time to obtain 20% harm reduction?

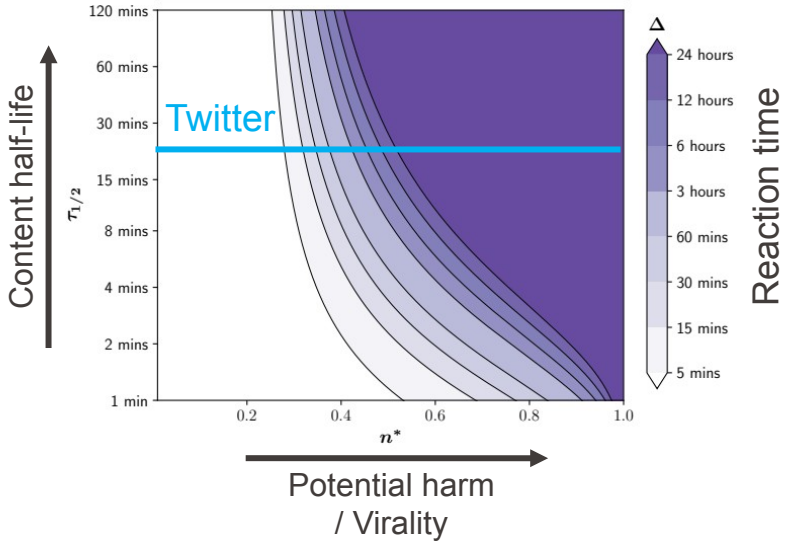

What is the reaction time to obtain 20% harm reduction?

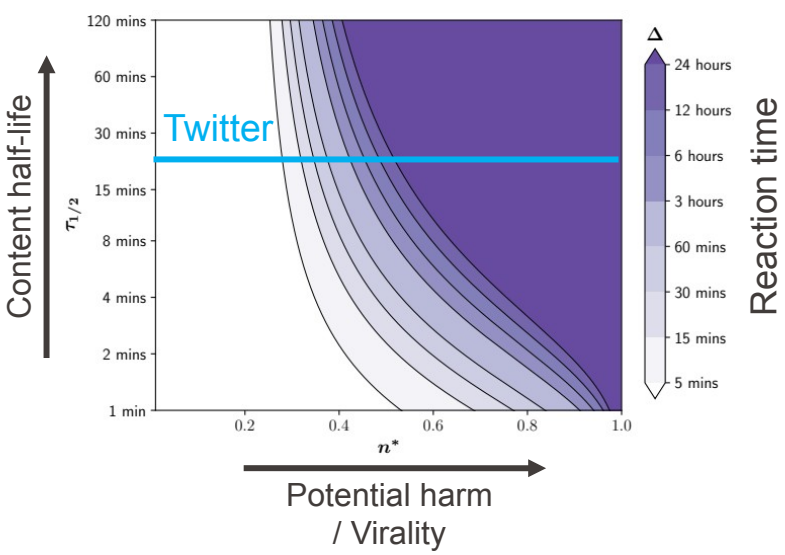

What is the achieved harm reduction when removing content after 24 hours?

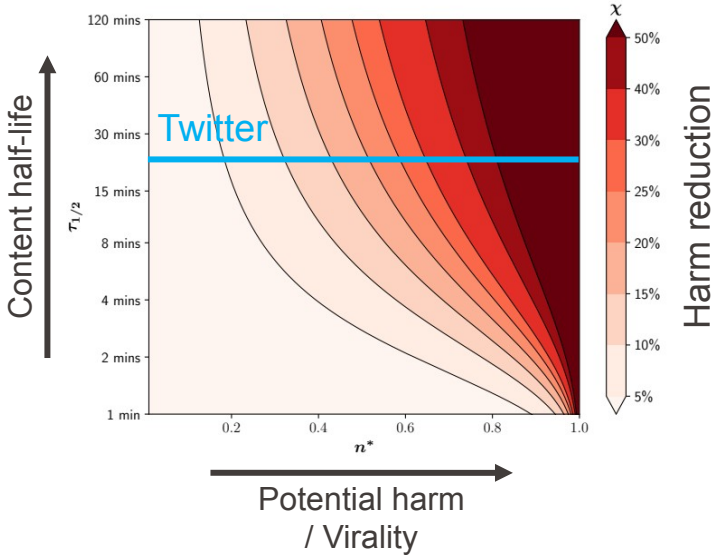

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

'On-the-fly' (real-time) parameter estimates

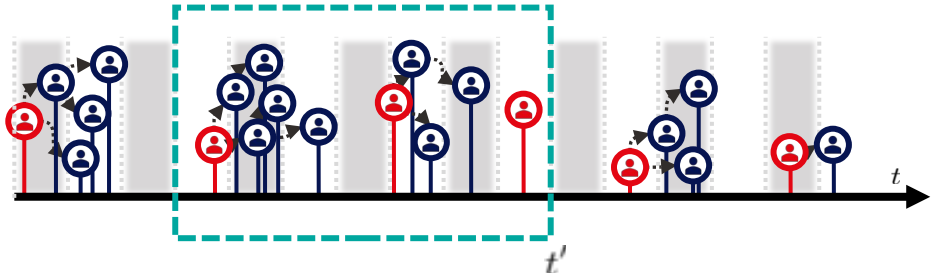

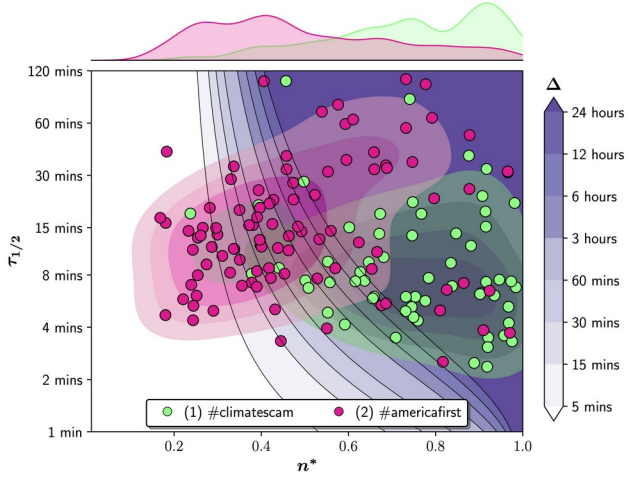

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

'On-the-fly' (real-time) parameter estimates

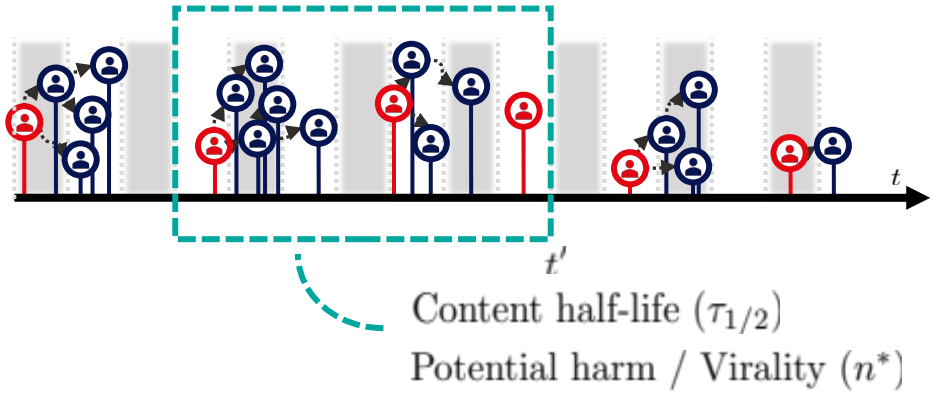

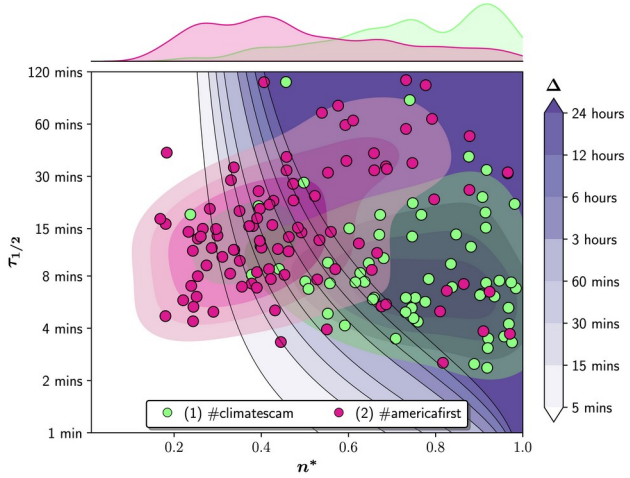

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

'On-the-fly' (real-time) parameter estimates

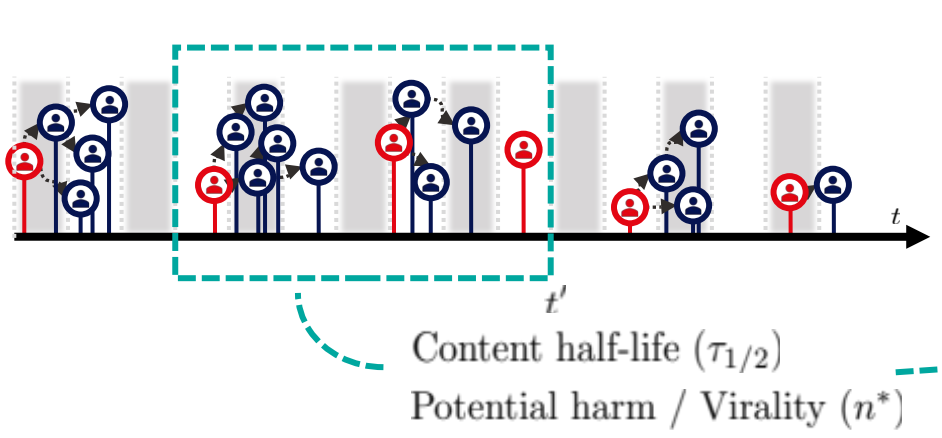

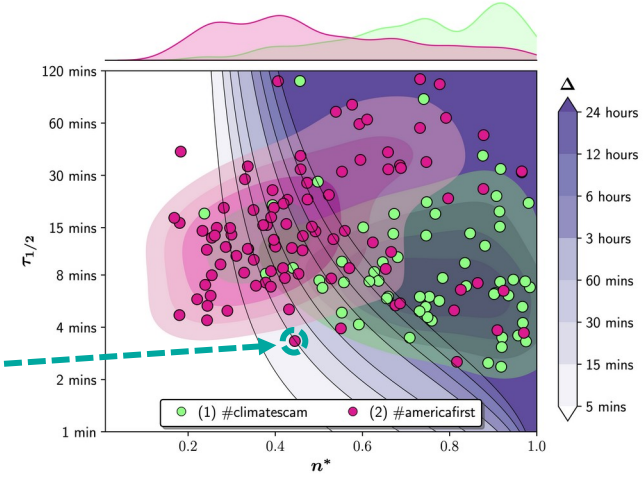

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

'On-the-fly' (real-time) parameter estimates

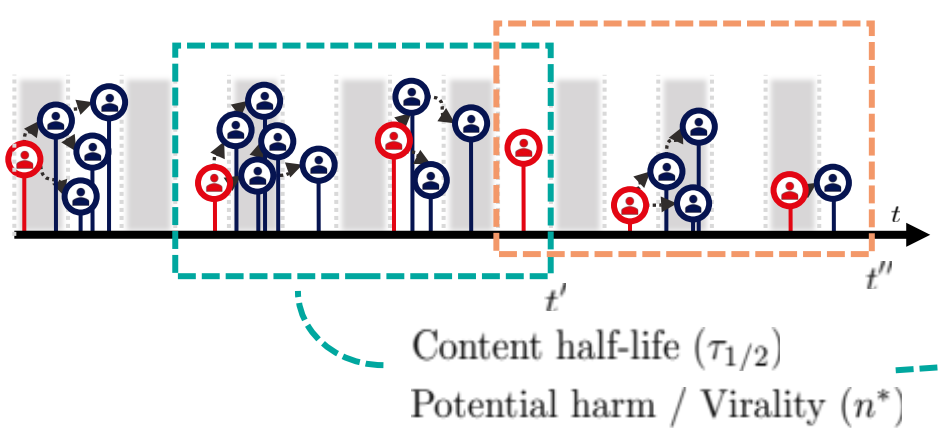

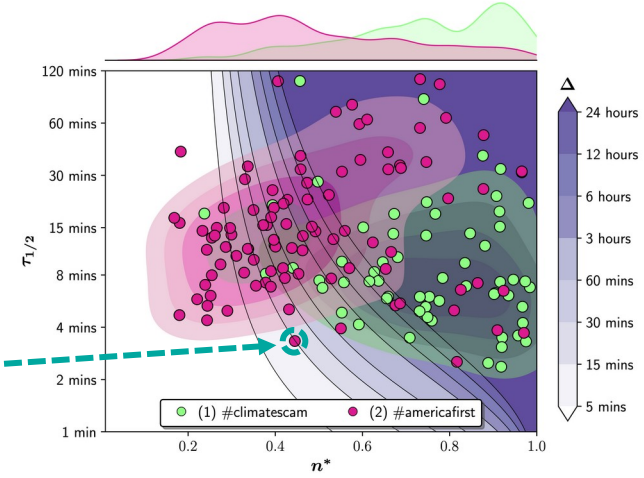

Application to real-world discussions

Twitter datasets (1 July to 31 December 2022)

- #climatescam (479,051 posts) – Controversial opinions regarding climate change [9]
- #americafirst or #americansfirst (278,899 posts) – Debates over key US political topics such as immigration and foreign policies [10]

'On-the-fly' (real-time) parameter estimates

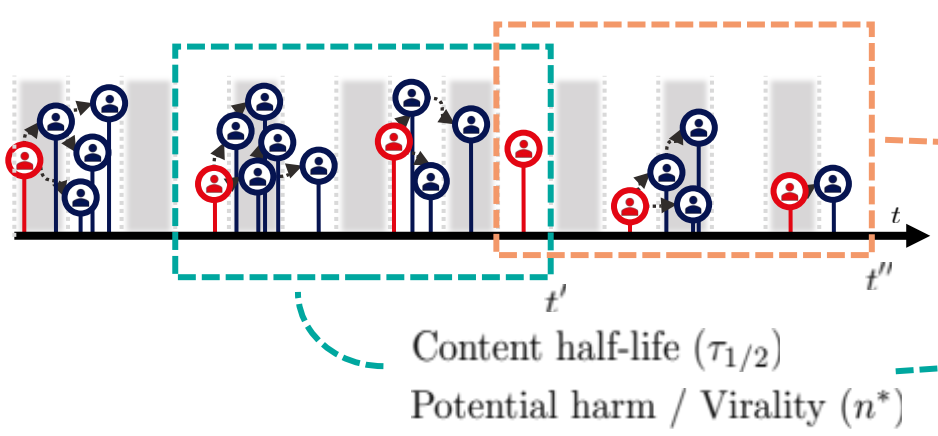

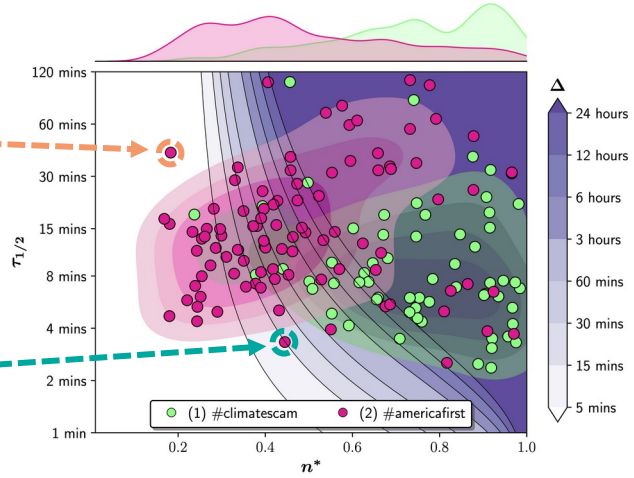

Effectiveness of EU-regulated moderation

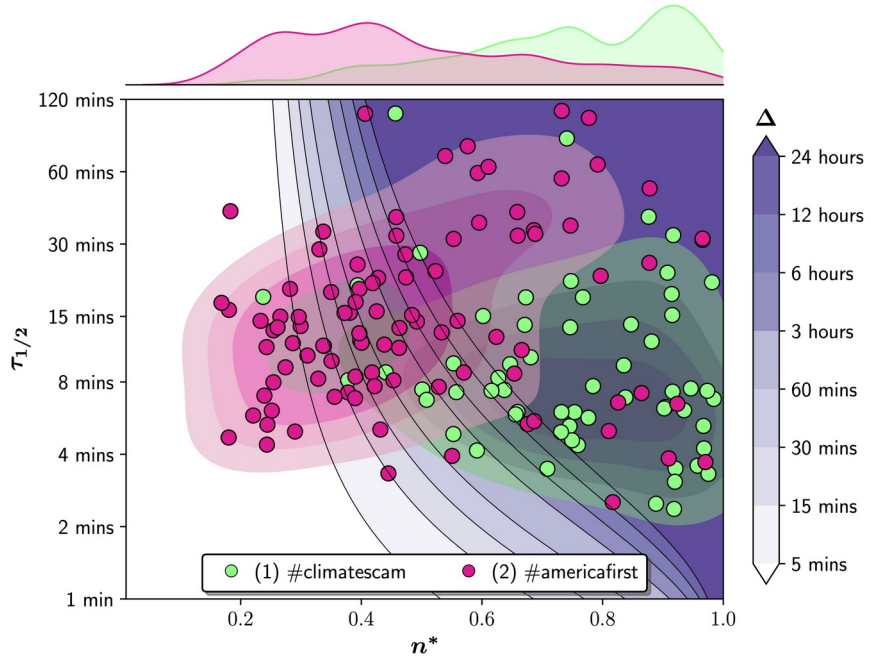

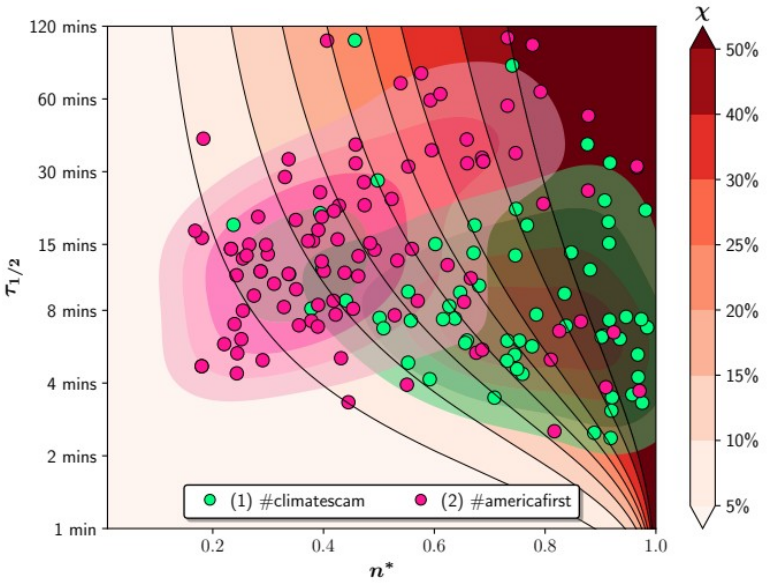

- Real-world potentially problematic content exhibits widely highly variable dynamics
- Harm reduction via manual flagging efforts is achievable

Effectiveness of EU-regulated moderation

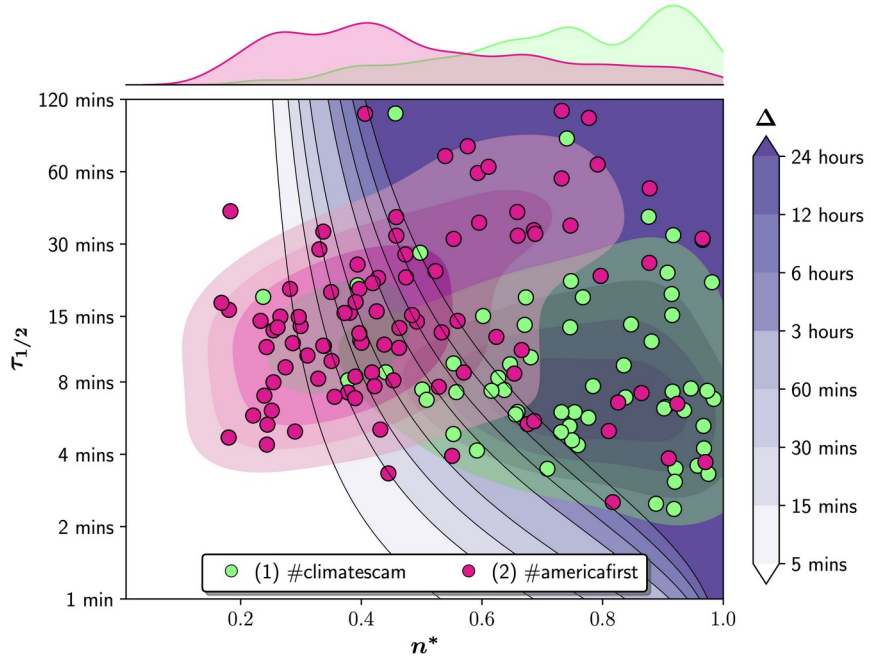

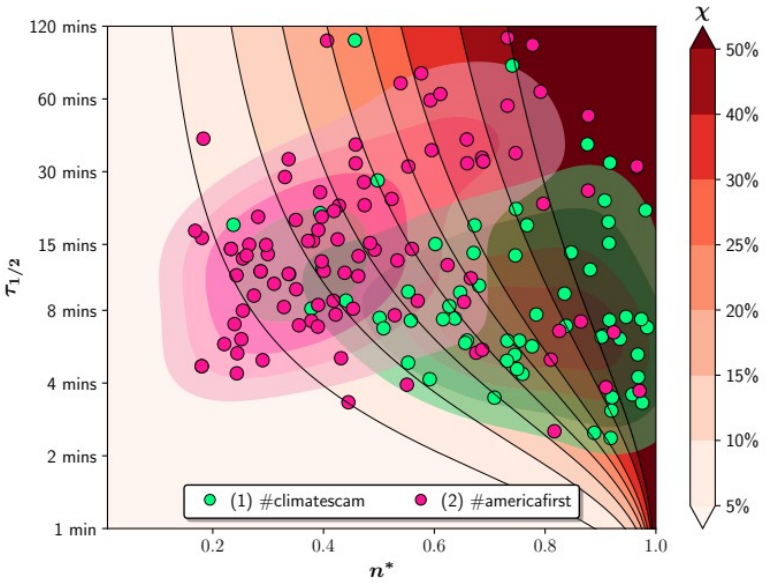

- Real-world potentially problematic content exhibits widely highly variable dynamics
- Harm reduction via manual flagging efforts is achievable

Topics	Potential harm / Virality	Content half-life	Harm reduction
▪ #climatescam	0.75	7.48 min	29.18%
▪ #americafirst	0.44	13.97 min	13.29%

Conclusion

- Harm reduction is achievable with manual flagging efforts, even for fast-paced platforms such as Twitter
- Framework for policymakers to draft mechanisms for content moderation by indicating where to focus human fact-checking efforts and how quickly to react

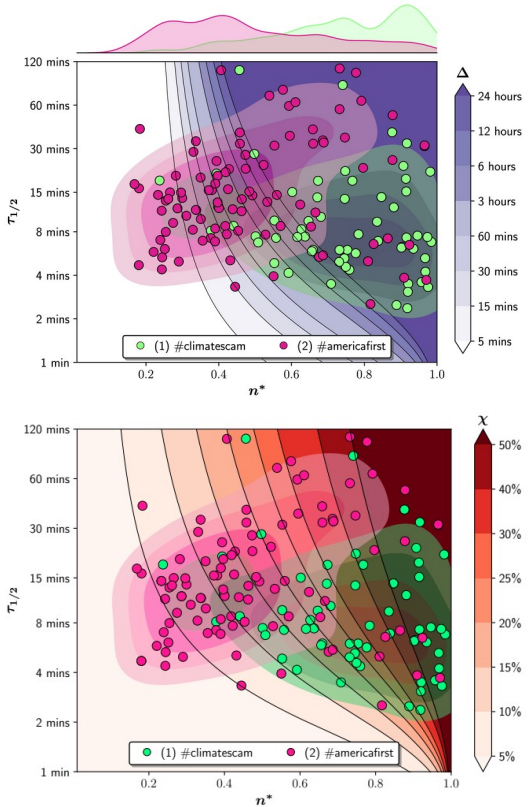

Conclusion

- Harm reduction is achievable with manual flagging efforts, even for fast-paced platforms such as Twitter
- Framework for policymakers to draft mechanisms for content moderation by indicating where to focus human fact-checking efforts and how quickly to react

Future work / Open questions

- How to select discussion topics?
- What is the ‘actual’ reaction time?
- What metrics measures the effectiveness in more granularity?
  - Post-based potential harm?

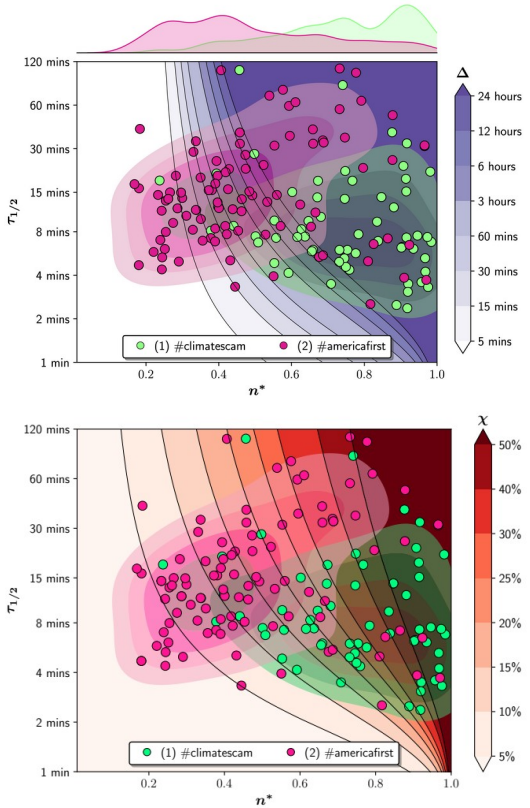

References

- [1] P. M. Barrett, *Who Moderates the Social Media Giants?* (NYU Stern Center for Business & HumanRights, 2020).
- [2] C. St. Aubin, J. Liedke, *Most Americans favor restrictions on false information, violent content online*. Pew Research Center (2023). https://www.pewresearch.org/short-reads/2023/07/20/most-americans-favor-restrictions-on-false-information-violent-content-online/ (Accessed 27 September 2023).
- [3] J. Singh, *OpenAI says AI tools can be effective in content moderation*. Reuters (2023). https://www.reuters.com/technology/openai-says-ai-tools-can-be-effective-content-moderation-2023-08-15/ (Accessed 27 September 2023).
- [4] J. J. McCorvey, *Tech layoffs hit "trust and safety" teams, raising fears of backsliding efforts to curb online abuse*. NBC News (2023). https://www.nbcnews.com/tech/tech-news/tech-layoffs-hit-trust-safety-teams-raising-fears-backsliding-efforts-rcna69111 (Accessed 27 September 2023).
- [5] R. Williams, *Humans may be more likely to believe disinformation generated by AI*. MIT Technology Review (2023). https://www.technologyreview.com/2023/06/28/1075683/humans-may-be-more-likely-to-believe-disinformation-generated-by-ai/ (Accessed 27 September 2023).
- [6] G. Spitale, N. Biller-Andorno, F. Germani, *AI Model GPT-3 (dis)informs us better than humans*. Science Advances, 9(26), eadh1850 (2023).
- [7] Regulation (EU) 2022/2065 of the European Parliament and of the Council of 19 October 2022 on a Single Market for Digital Services and amending Directive 2000/31/EC (Digital Services Act).OJ L277, 1–102 (2022).
- [8] S. M. Graffius, *Lifespan (half-life) of social media posts: Update for 2023* (2023). https://dx.doi.org/10.13140/RG.2.2.19783.98722 (Accessed 31 July 2023).
- [9] O. Milman, *#ClimateScam: Denialism claims flooding Twitter have scientists worried*. Guardian (2022). https://bit.ly/guardian- climatescam-twitter (Accessed 31 July 2023)
- [10] D. L. Linvill, P. L. Warren, *Troll factories: Manufacturing specialized disinformation on Twitter*. Polit. Commun. 37, 447–467 (2020).

Applicability to time-censored information

Base intensity Self-excitation

Intensity function $\lambda(t|\mathcal{H}_t) = \mu + \sum_{i:t_i<t} \phi(t-t_i)$

Kernel function

Event history

$\mathcal{H}_T = \{t_1, \ldots, t_K\} \subset (0, T]$

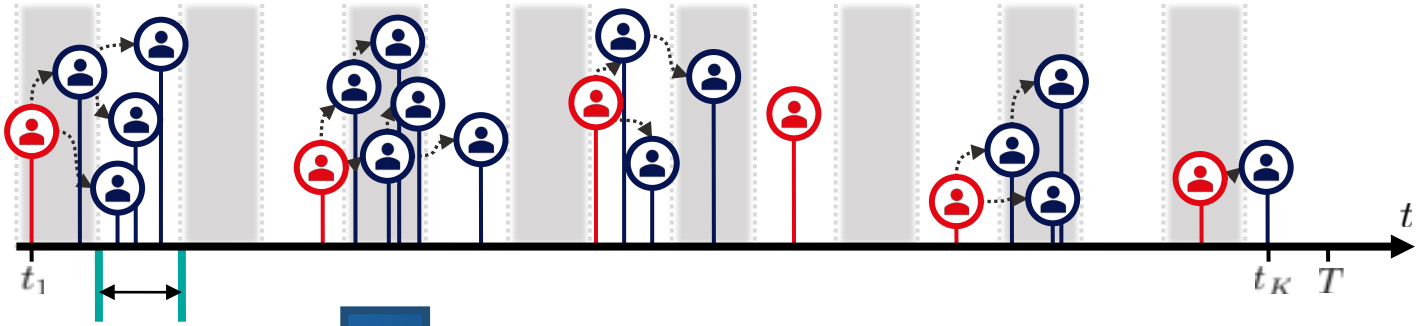

Bin-count vector

$X = (X_1, \ldots, X_L)$

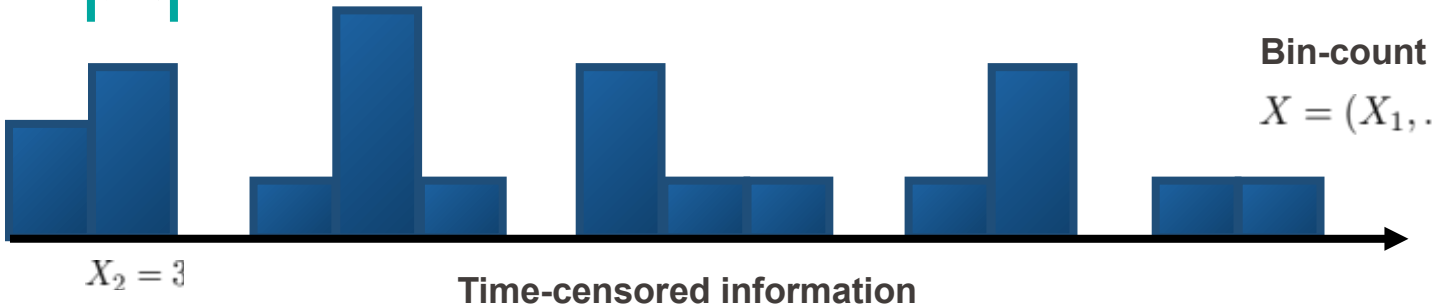

Time-censored information